



**High-Performance Data-Intensive
Computing Systems Laboratory**

Scaling Research with the NSF Nautilus HyperCluster

A Tutorial and Case Study



University of Missouri

November 4, 2022

Overview

► Introduction

- Docker
 - Containerization
 - Key Concepts: Images, Containers, Registries
- Kubernetes
 - Container Orchestration
 - Key Concepts: Nodes, Namespaces, Pods, Jobs

► NSF Nautilus HyperCluster

- Resources Available
- System Overview
- By the Numbers

► Using Nautilus: Tutorial

- Getting Started
- Using Persistent Storage
 - S3 Cloud Storage Integration
- Creating a Pod
- Creating a Batch Job
- Using Nautilus for Deep Learning
- Automating Deep Learning Job Deployments


Docker

► What is Docker?

- Docker is a set of platform as a service products that use OS-level virtualization to deliver software in packages called containers.¹
- You can think of Docker containers as mini-VMs that contain all the packages, both at the OS and language-specific level, necessary to run your software.

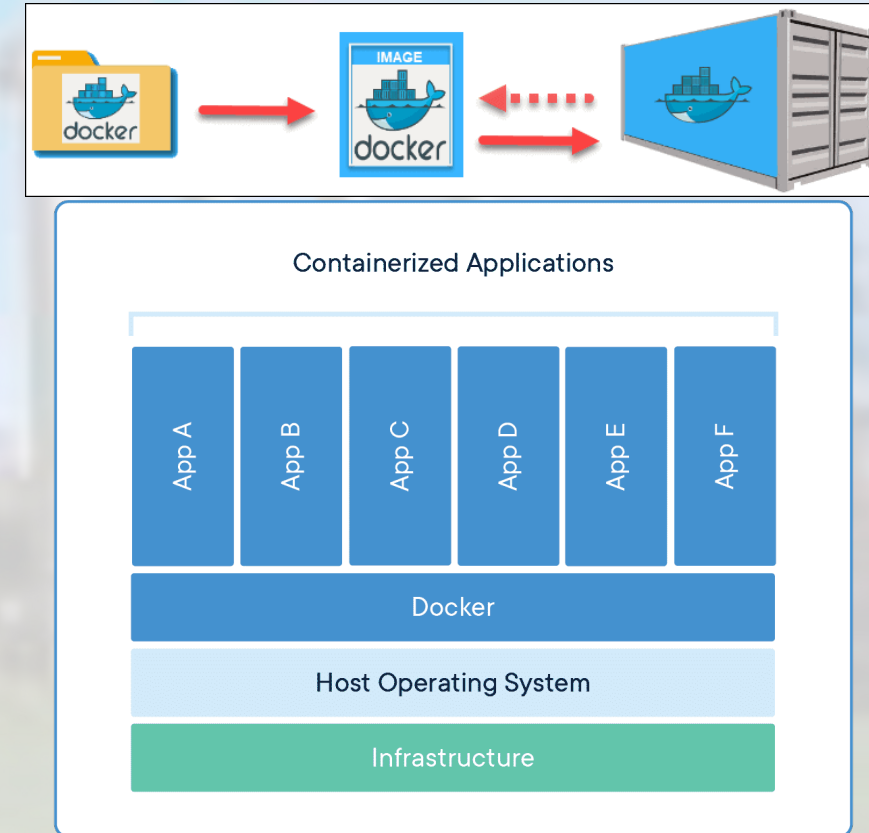
► Why Docker?

- Docker enables predictable and reliable deployment of software.
- Docker containers are portable!
 - local development computers, compute clusters, internal compute servers, cloud infrastructure, and more!
- *Docker containers are how software is deployed in **Kubernetes**!*



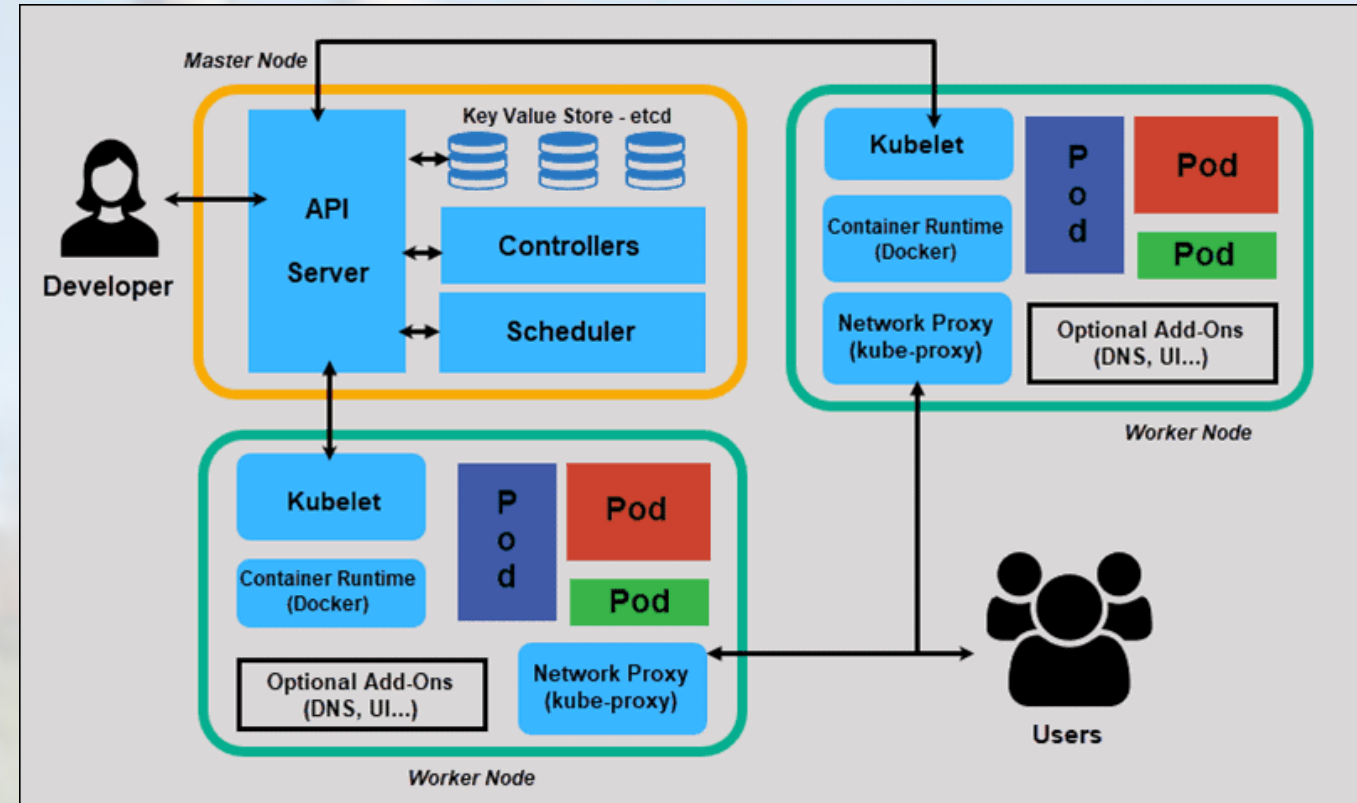
Key Docker Concepts

- ▶ **Dockerfile** - A list of commands and instructions describing how to build an Image
- ▶ **Image** - Standard unit of software that packages up code and its dependencies so the application runs reliably from one computing environment to another.
 - ▶ Includes everything needed to run an application: code, runtime, system tools, system libraries and settings.
- ▶ **Container** - An instantiated runtime of a docker image, containing all necessary software for a given application to run, both at the OS and language-specific level
 - ▶ Images become containers at runtime
- ▶ **Registry** - a service for storing private container images



Kubernetes


- Kubernetes, also known as K8s, is an open-source system for automating deployment, scaling, and management of containerized applications.¹
- Kubernetes enables container orchestration and serves as the backbone of the NSF Nautilus HyperCluster



Key Kubernetes Concepts




- ▶ **Nodes** - a node is a physical machine where containers are deployed. Each node must run a container runtime (in Nautilus, it is 'Containerd').
- ▶ **Namespaces** - provide a way for K8s to partition cluster resources across multiple or many users in an exclusive way.
- ▶ **Pods** - pods are the basic scheduling unit of K8s. Pods consist of one or more containers running inside. Each pod has a unique IP address to enable micro services or applications
- ▶ **Jobs** - long running processes that require more than 6 hours of runtime or have more demanding compute that 2 CPU cores, 8 GB or RAM, and 1 GPU
- ▶ **Services** - a set of pods working together

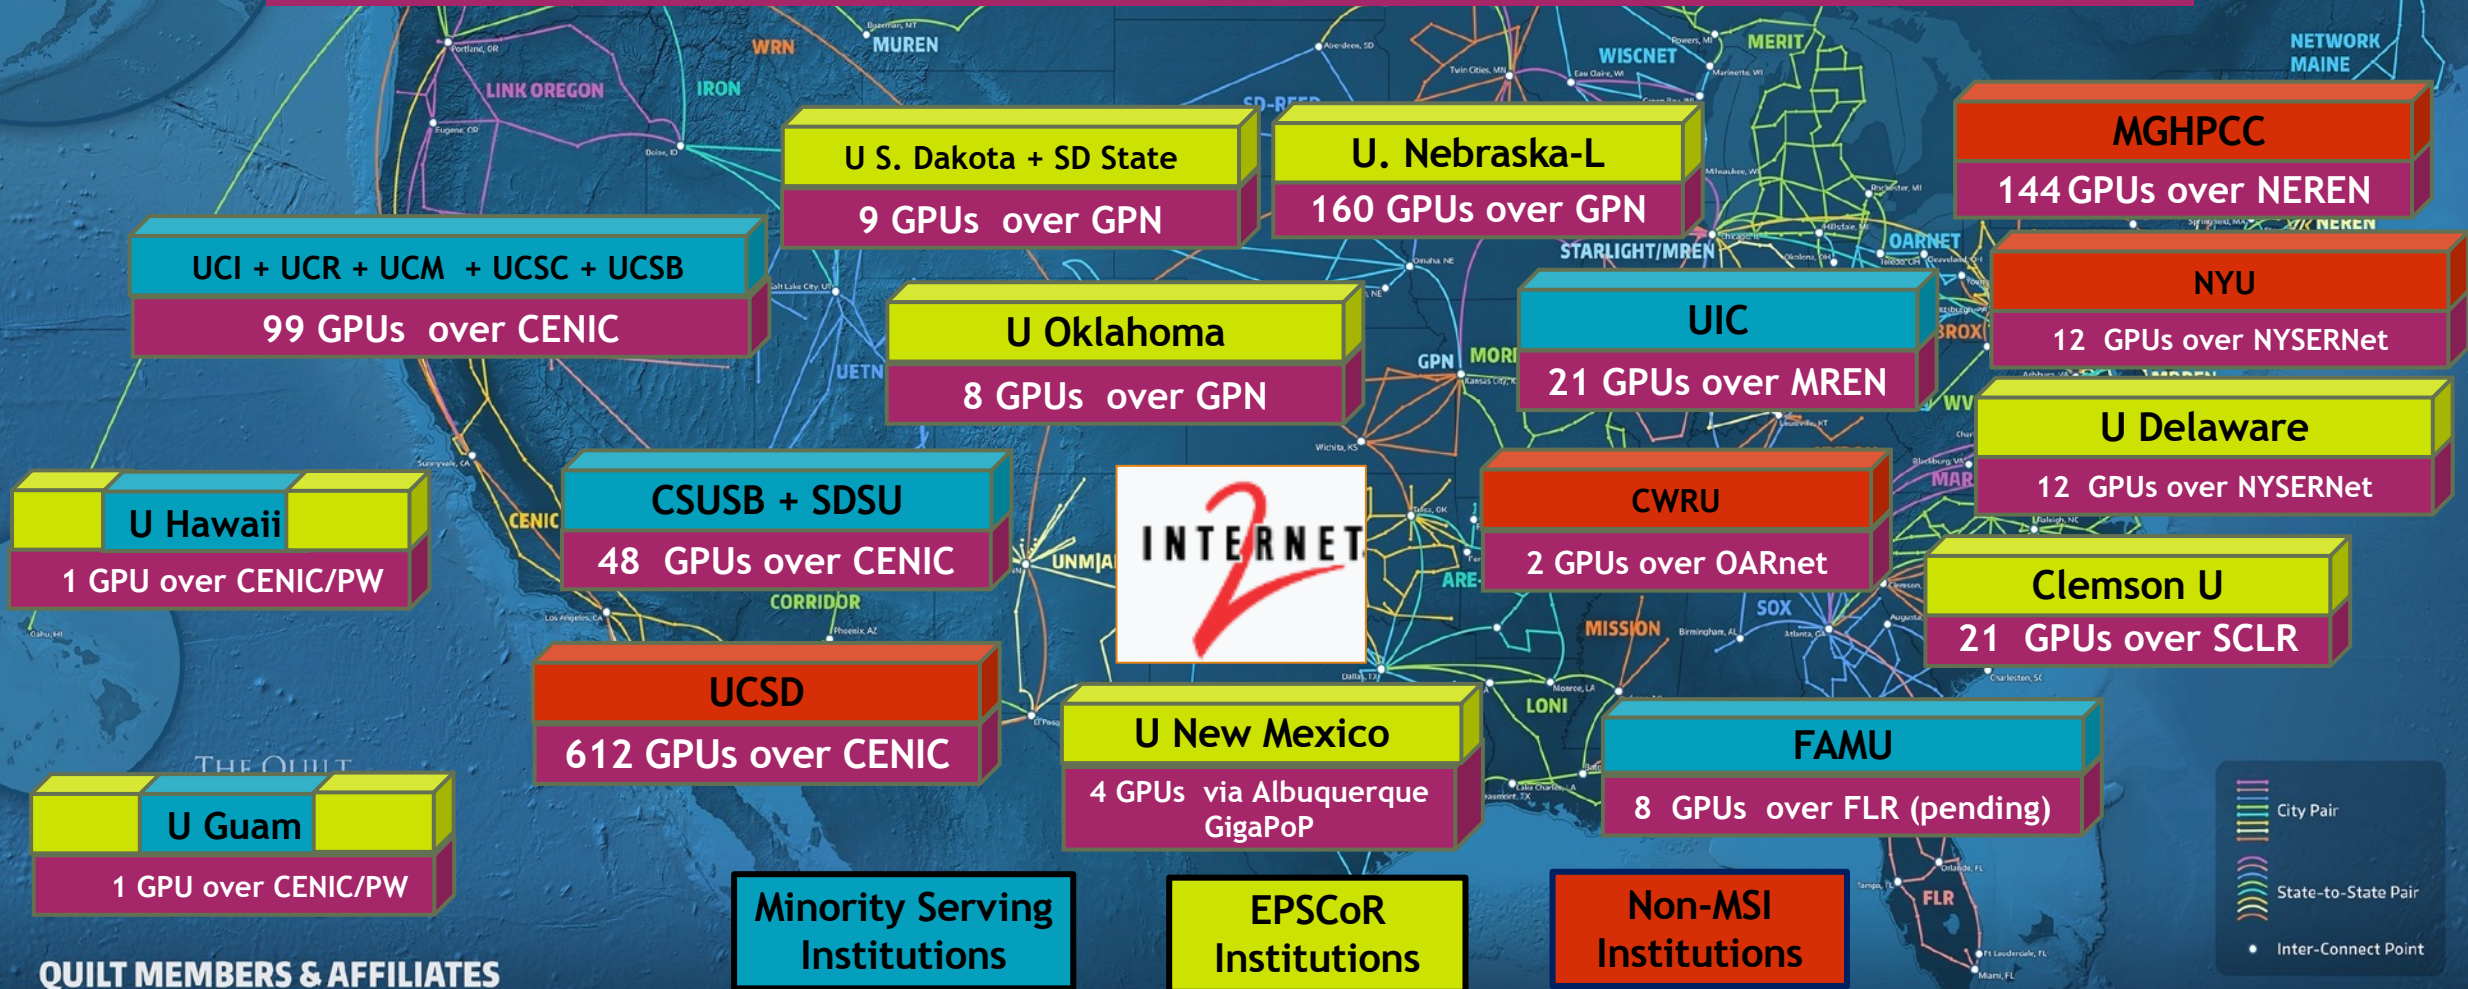


NSF Nautilus HyperCluster

- ▶ The NSF Nautilus HyperCluster is a Kubernetes cluster with vast resources that can be utilized for various research purposes:
 - ▶ Prototyping research code
 - ▶ S3 cloud storage for data and models
 - ▶ Accelerated small-scale research compute
 - ▶ Scaling research compute for large scale experimentation
- ▶ Resources readily available:
 - ▶ Nodes: 263
 - ▶ Logical CPU Cores: 2,114
 - ▶ RAM: 12.47 TB
 - ▶ GPUs: 158



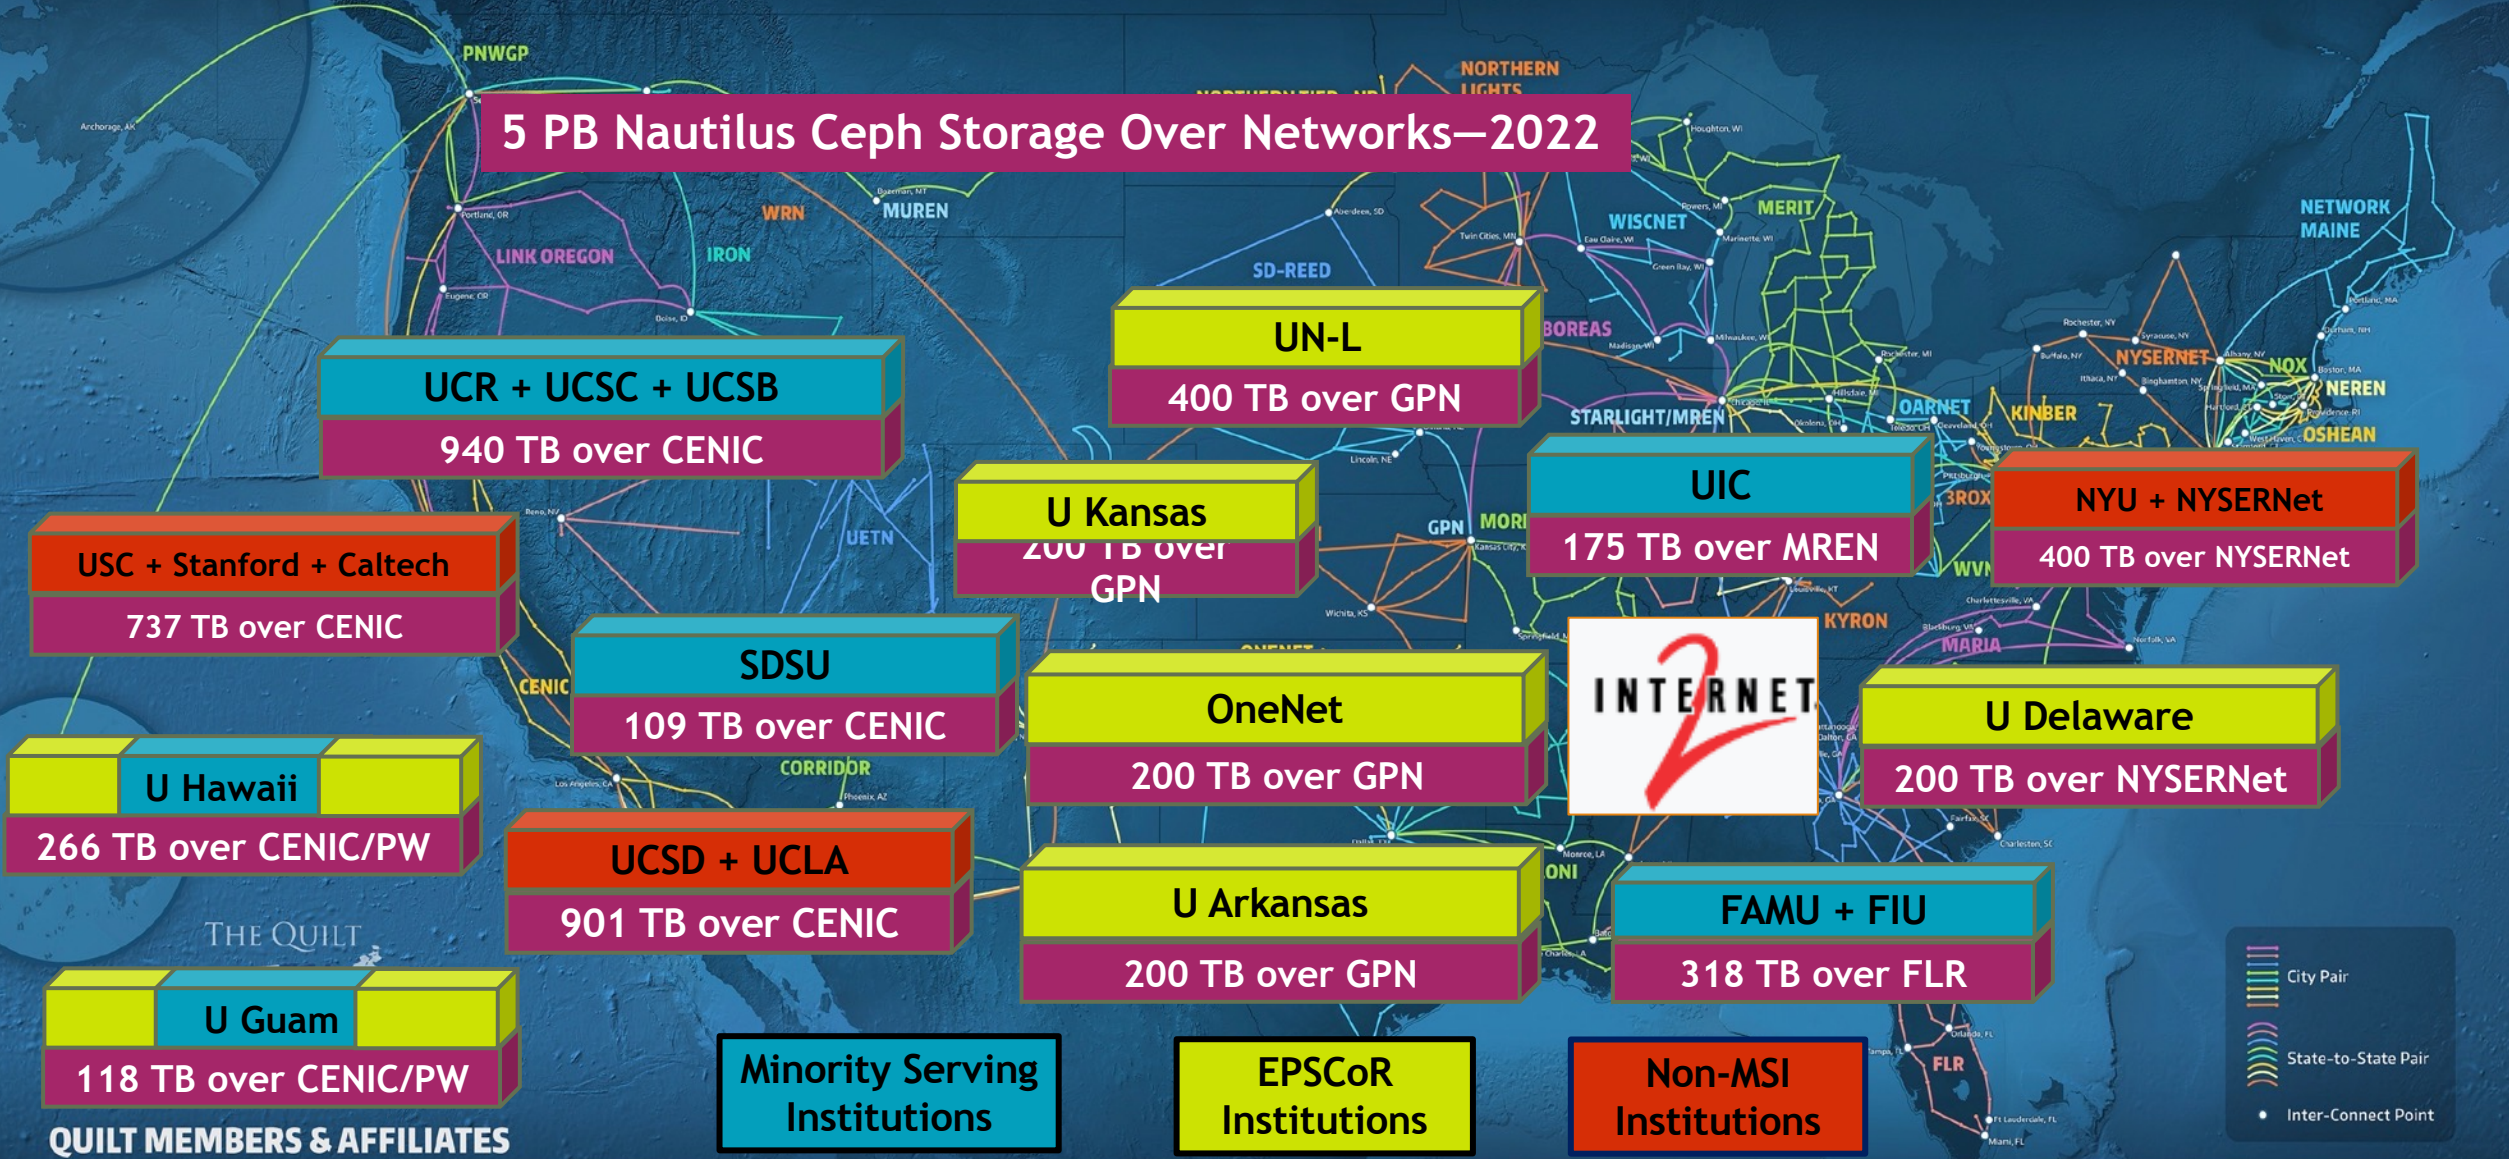
Nautilus ~1,100 GPUs Distributed over US Networks—Fall 2022



QUILT MEMBERS & AFFILIATES



5 PB Nautilus Ceph Storage Over Networks—2022



Container Orchestration in Nautilus

1. Users push code to VCS
 - <https://gitlab.nrp-nautilus.io/>
2. GitLab CI/CD build Docker image and push to Container Registry
3. Users create pod/job using published Docker image
 - KubeCTL Command Line Tool
4. Kubernetes schedules the pod/job onto a node with required resources





Figure 4: Model Container Supply Chain

Case Study System Overview:


Nautilus for Accelerated Deep Learning Research

- ▶ Using containerized model definition and list of jobs
- ▶ Mounted persistent data storage to each pod
- ▶ Each GPU job produces an associated trained model
- ▶ Automation currently performed via environment variables and bash, but more sophisticated methods in development
- ▶ Models are later sync'd to Nautilus S3 bucket for later use in evaluation or other ML applications






Deep Learning on Nautilus: By the Numbers




9 Deep Neural
Architectures




27 Trained Models




124.74 TB
Imagery
Processed



3 Open Source
HR-RSI Datasets




- 8,100 Epochs
- 30M iterations
- 1.7B parameters




Wall Clock Time:
76 days, 10 hours

Non-California Nautilus PI Namespace 2021 Usage by State: “Big MO!”

Data/Plots provided by Larry Smarr (PI, National Research Platform & father of US Super Computing Centers)



17,217 GPU-hrs



28,088 CPU core-hrs



University of Missouri

Grant Scott, UMC
Helped Organize the UMC PRP Usage